

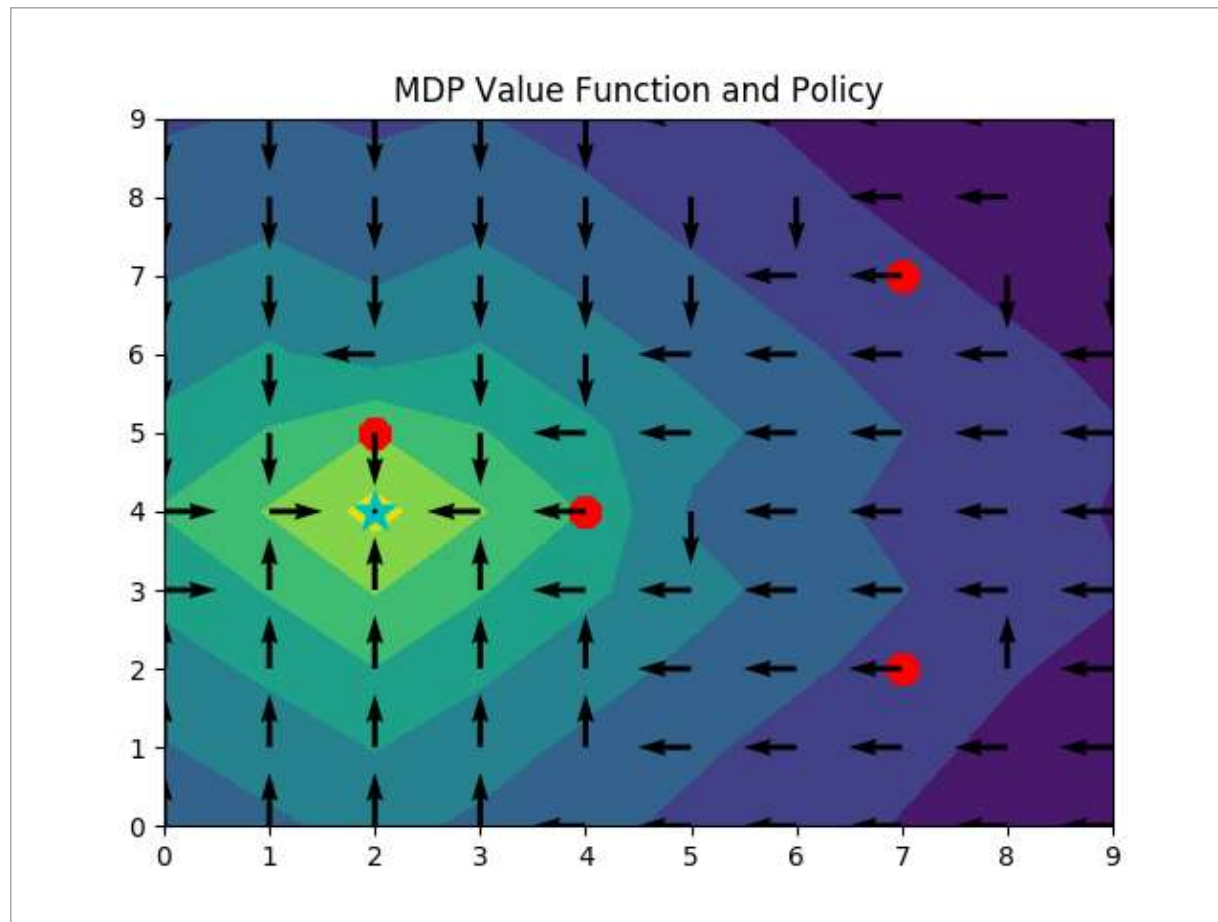
ASEN 6519, Spring 2017

Algorithms for Aerospace Autonomy

Lecture 24: Approximate POMDPS


Luke Burks
04/20/2017

MDP: Modified Europa Hopper Problem





MDP: The simpler hallway problem

- A robot is in a hallway trying to find the power outlet
- Also, there's lava at each end of the hallway...

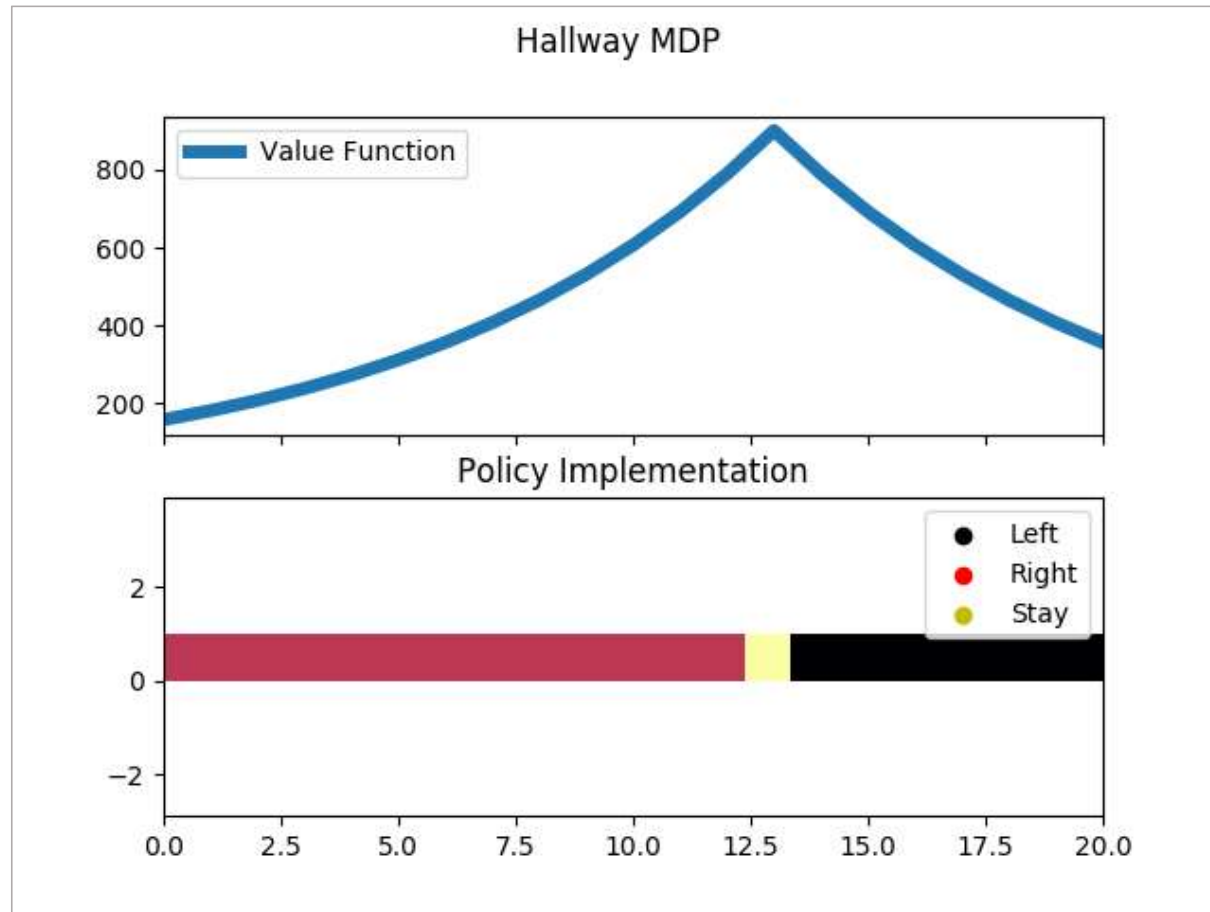


MDP: The simpler hallway problem

- The hallway is discretized into 21 spaces
- The robot can take 3 actions, and receive 5 observations



MDP: The simpler hallway problem



Why this is hard with a POMDP

- In the (super weird) baby problem from the book there were two states, so we could visualize the alpha vectors easily enough.
- Here the value function is defined by a 21 dimensional piece-wise set of alpha vectors...
- That awful set of calculations we did on Thursday, the additional actions and observations lead to far more of those
 - 2 step look ahead: 729 alphas
 - 4 step look ahead: 10^{74} alphas
 - 5 step look ahead: 10^{372} alphas
 - 20 step look ahead: $10^{10^{13}}$ alphas

QMDP: What if we pretend it's actually an MDP?

- QMDP [Littman et. Al 1995]
- Calculates the MDP-Value of taking each possible action


$$Q(s_i, a) = r(s_i, a) + \sum_j^N V_{MDP}(s_j) p(s_j | s_i, a)$$

- The policy for a belief is the action that maximizes the expectation of that value

$$\pi(b) = \operatorname{argmax}_a \sum_i^N b(s_i) Q(s_i, a)$$

- Is exact for cases where the state becomes fully observed after the action
- Can be shown to be an upper bound on the value function
- **Preforms poorly at pure information gathering actions**

QMDP Simulation Example




FIB: What if we only kind of pretend its actually an MDP?

- Fast Informed Bound [[Hauskrecht 2000](#)]
- Takes all possible observations into account when calculating the value of each action



$$Q(s_i, a) = r(s_i, a) + \sum_j^N \sum_o^{|O|} p(o|s_j) p(s_j|s_i, a) V_{MDP}(s_j)$$

- The value estimate is upper bounded by QMDP and lower bounded by the true value
- Takes longer than QMDP, especially for problems with lots of observations

MDP: Just a reminder




QMDP vs. FIB




PBVI: Do we have to care about ALL the beliefs?

- Point-Based Value Iteration [[Pineau et. al 2003](#)]
- Given some starting belief, there are many beliefs you'll never encounter no matter the sequence of actions and observations
- Use what you know about the beliefs you've seen to tell you what to do at new beliefs



PBVI: Do we have to care about ALL the beliefs?



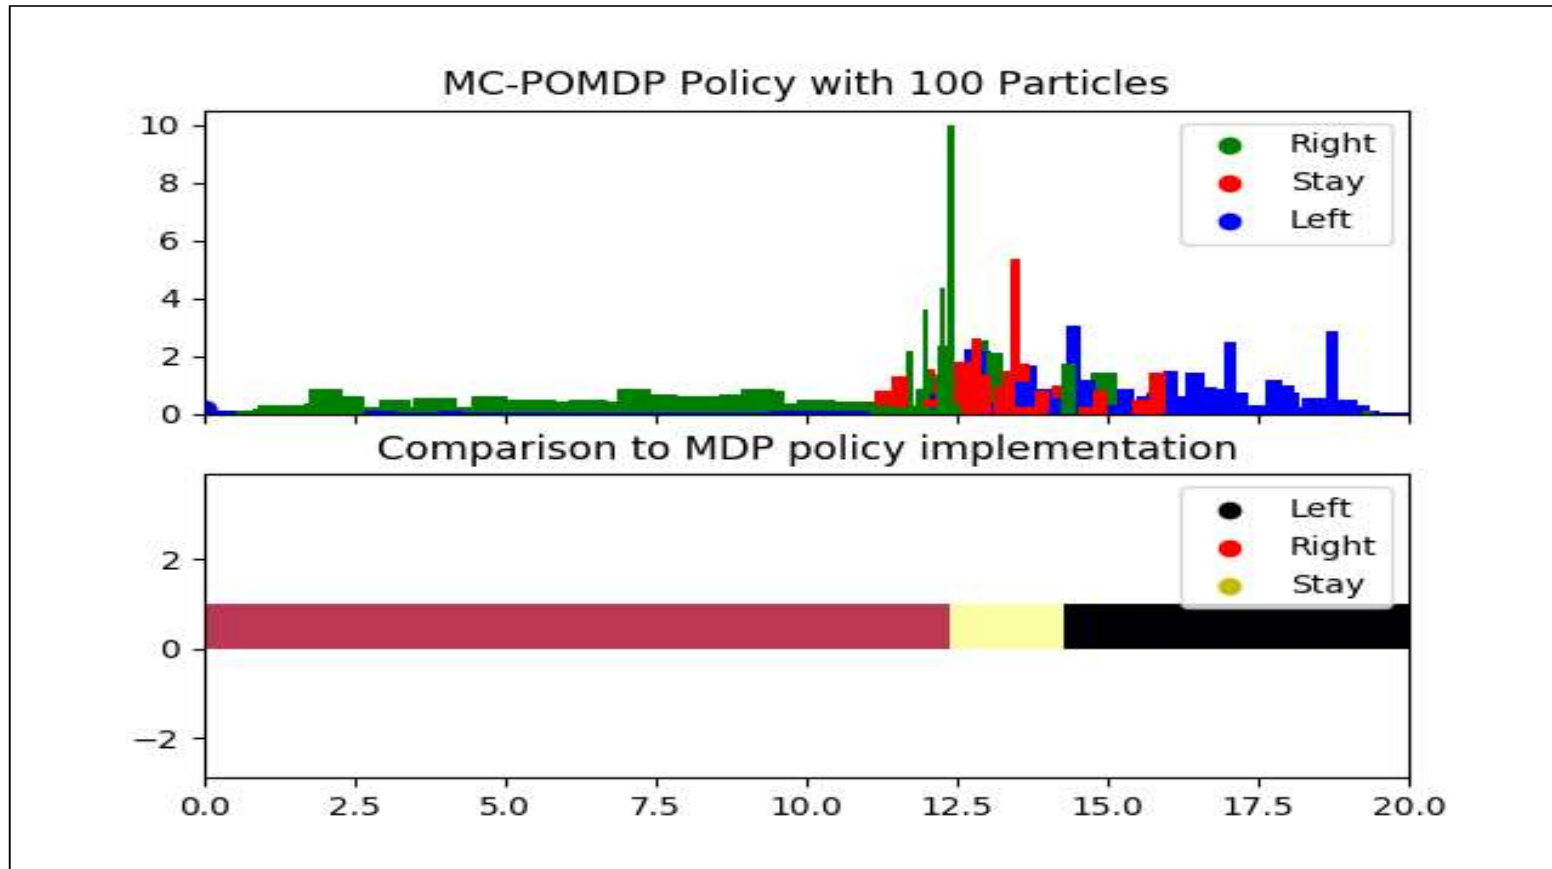
MC-POMDP: Because putting Monte Carlo in front of anything makes it possible right?

- Monte Carlo POMDP [Thrun 2003]
- Represent the belief by a set of (x) particles

$$b(s) \rightarrow \mathcal{X}, x \in \mathcal{X}$$

- Push those particles through your transition function
- Use importance resampling to kick out particles that end up in unlikely places
- Each particle set you find has a value calculated for it
- Can use nearest neighbor interpolation between particle sets to find the action for a new set
- Works in continuous spaces

MC-POMDP: Because putting Monte Carlo in front of anything makes it possible right?




POMCP: When your problem requires a state for every atom in the solar system...

- Partially Observable Monte Carlo Planning [\[Silver 2010\]](#)
- Online POMDP Solver
- Plans forward from the current belief state
 - Guarantees reachable beliefs are explored
- Cares about histories, not beliefs
 - Histories are sequences of past actions and observations
 - Given a prior belief and a history, you can always construct the posterior belief
- For history h_t using K particles,

$$\hat{\mathcal{X}}(s, h_t) = \sum_i^K \delta_{s\mathcal{X}_t^i}, \text{ where } \delta_{ss'} \text{ is the kronecker delta}$$

POMCP: When your problem requires a state for every atom in the solar system...

- Builds a tree of possible future histories
- When an action is taken and an observation received, the root is moved down the corresponding branches while other branches are pruned



POMCP: In action

- 10^{56} States
- 1024 Obs
- 4 Actions

